

CPU Cache

Contents

1. General definitions
 - 1.1. Cache
 - 1.2. Cache hit & miss
2. CPU cache
 - 2.1. CPU cache hierarchy
 - 2.2. CPU cache level characteristics
 - 2.3. Cache entry structure
3. Cache mapping
 - 3.1. Direct-mapped
 - 3.2. Fully associative
 - 3.3. Set-associative
4. Cache policies
 - 4.1. Writing policies
 - 4.2. Replacement policies
5. Cache coherence
6. Going further

1.1- General definitions: cache

- Hardware or software component.
- Stores the result of expensive operations or computations:
 - I/O operations : opened file cached in RAM to limit hard drive operations,
 - Expensive function calls : lookup tables using parameters = memoization,
 - Memory access,
 - ...
- Accelerate access to that data.

1.2- General definitions: cache hit & miss


- Cache hit: request for data that was stored in the cache.
- Cache miss: request for data that was not stored in the cache.
 - Compulsory miss: first ever request for the data, cache miss is inevitable !
 - Capacity miss: the cache is not large enough to hold the working set
 - Conflict miss: the requested data was previously evicted from the cache
 - Coherence miss: the requested data was invalidated by another thread
 - Coverage miss: the requested data was invalidated by the coherence directory

2- CPU cache

- Hardware component
- Used since the 1980s
- Lies between the main memory and the CPU registers
- Bridges operating frequency gap between CPU and main memory
- Uses static rather than dynamic RAM (SRAM/DRAM)
 - Expensive
 - Low power consumption
 - No refresh needed
- Benefits from locality of reference (both spatial and temporal)

2.1- CPU cache hierarchy


- Modern CPUs have multiple level of cache
 - Most commonly 2 or 3
- Level 1 cache is split:
 - L1d -> data cache
 - L1i -> instruction cache



2.2- CPU cache level characteristics


Each cache level can be:

- Shared or Local
- Separate or Unified
- Inclusive,
 - L_{n+1} contains data stored in L_n
 - Remove data by checking high level caches only
 - Echo cache line eviction to lower level caches
- Exclusive,
 - L_{n+1} does NOT contain data stored in L_n
 - Stores more data
 - Both cache must have same line size
- or neither
 - L_{n+1} may or may not contain data stored in L_n



2.3- Cache entry structure

- Data is transferred between the main memory and the cache in fixed size blocks called *cache lines* or *cache blocks*
 - Size between 32 and 256 bytes, most commonly 64 bytes
 - Size may vary between cache levels, higher size reduce compulsory miss rate (spatial locality)
- When a cache line is copied from the main memory to the cache, a cache entry is created. It holds:
 - The copied data,
 - A tag that identifies the entry,
 - Flag bits describing the status of the cache entry (validity, dirtiness ...)



- Only the size of the data counts toward the cache total size
 - A 256 KiB cache holds 256KiB of data !

3- Cache mapping

When we copy data from the main memory, where should it go ?

- 1 possible destination -> direct-mapped cache
 - Simplest design
 - More conflict misses and cache thrashing
- Every possible destination -> fully associative cache
 - Costly with large caches: must check EVERY cache block
 - No conflict misses, only capacity
- A compromise between the two -> set-associative cache
 - A main memory block has 2 to 16 possible destinations
 - Fewer conflict misses

3.1- Direct-mapped cache

Example:

- Main memory:
 - 256B (addresses: 0 to 255)
- Cache:
 - 32B direct-mapped
 - 8 cache lines
 - 4B line size
 - 3bits tag

	Tag	Data	Valid Bit
0	-	-	0
1	-	-	0
2	-	-	0
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.1- Direct-mapped cache: loading memory

- CPU request data from main memory with address 155.

155 \Leftrightarrow 10011011

- Block size is 4B = 2²B, set 2 least significant bits of address to 0 to find block address

10011011 & 11111100 = 10011000 \Leftrightarrow 152

- Right shift block address by 2 to find block index in main memory (equ. to division by 4)

10011000 >> 2 = 00100110 \Leftrightarrow 38

Example:

- Main memory: 256B (addresses: 0 to 255)

address	151	152	153	154	155	156
data	5 B	7 1	0 F	8 4	E 3	9 2

memory block

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
1	-	-	0
2	-	-	0
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.1- Direct-mapped cache: loading memory

- There are $8 = 2^3$ cache lines, keep the **3** least significant bits of that index to find the cache entry index (equ. to mod 8)

$00100110 \& 00000111 = 00000110 \Leftrightarrow 6$

- Use the **remaining** bits as the tag.

$(00100110 \& 11111000) \gg 3 = 00000100 \Leftrightarrow 4$

Example:

- Main memory: 256B (addresses: 0 to 255)

address	151	152	153	154	155	156
data	5 B	7 1	0 F	8 4	E 3	9 2

memory block

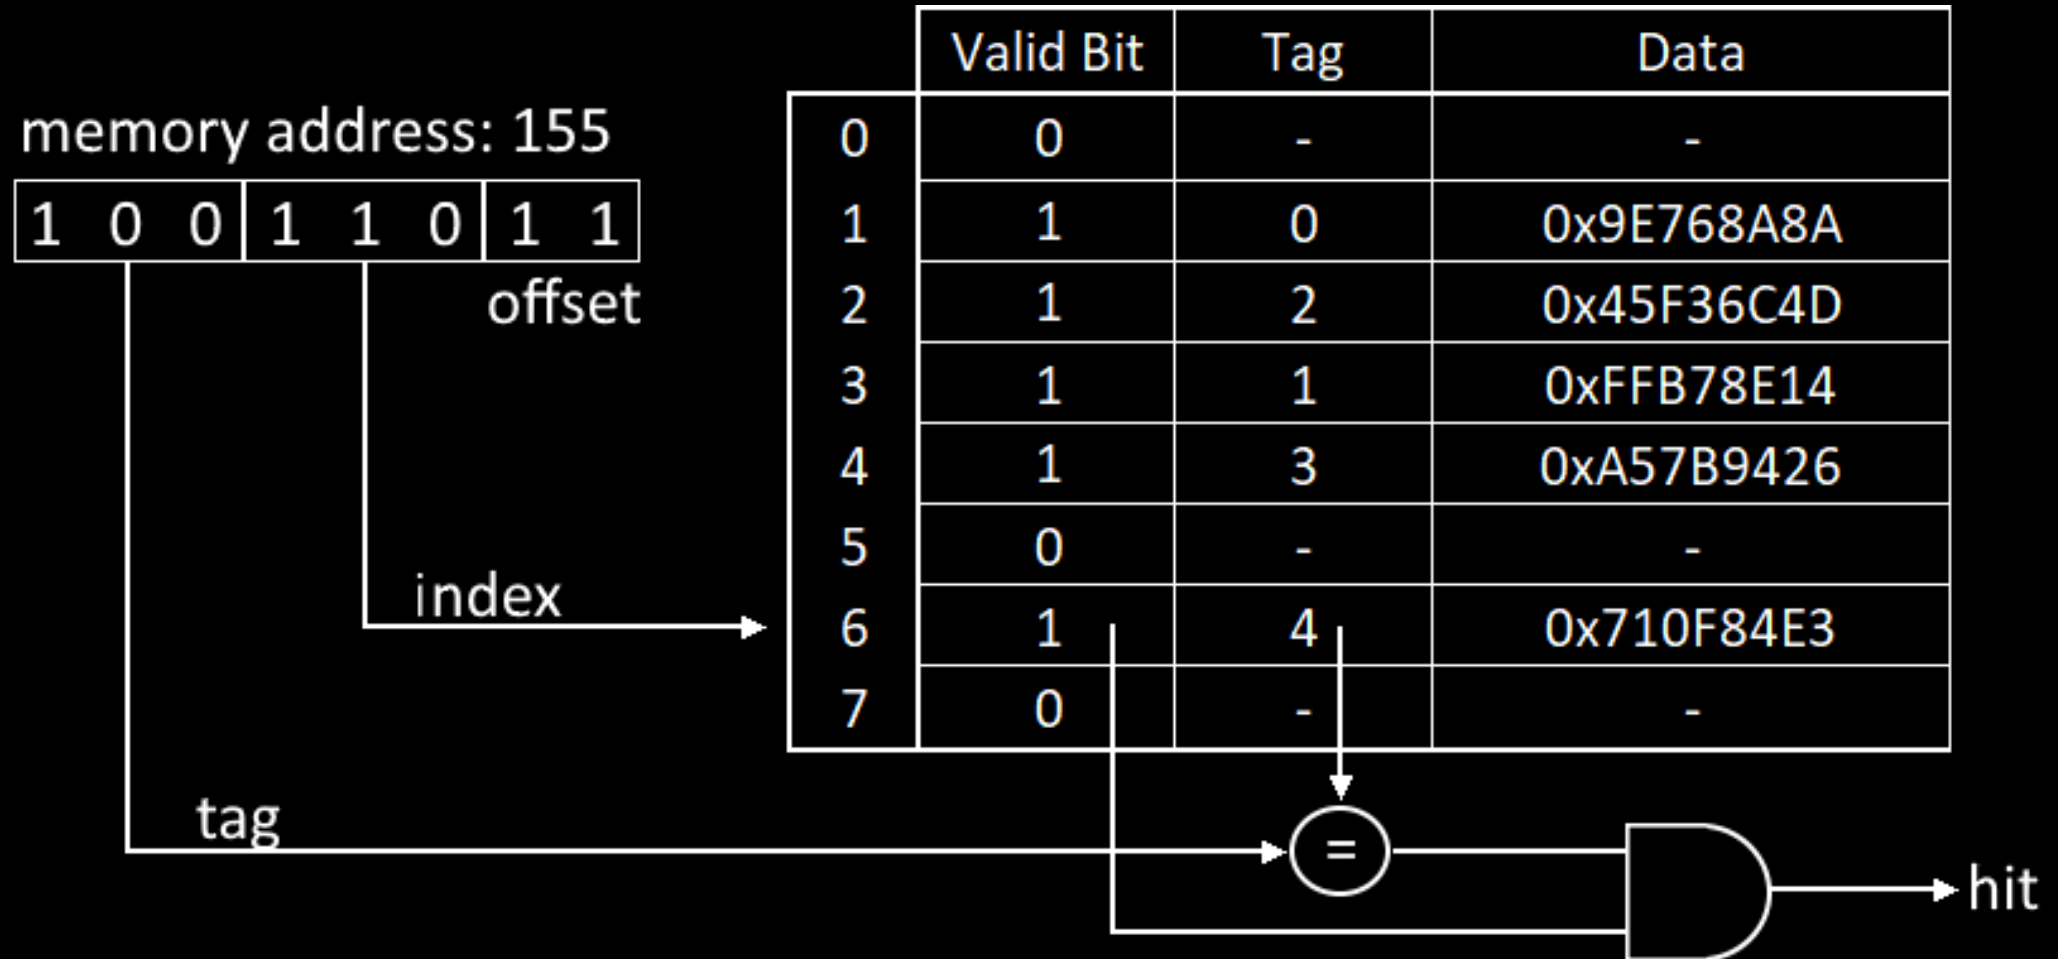
- Cache: 32B direct-mapped, $8 * 4B$ cache lines

	Tag	Data	Valid Bit
0	-	-	0
1	-	-	0
2	-	-	0
3	-	-	0
4	-	-	0
5	-	-	0
6	4	0x710F84E3	1
7	-	-	0

3.1- Direct-mapped cache: data request

Example:

- Main memory: 256B (addresses: 0 to 255)
- Cache: 32B direct-mapped, 8 * 4B cache lines



3.1- Direct-mapped cache: conflict misses and thrashing

- Conflict miss rate is high
- Very high conflict miss rate = thrashing

Example:

- Load address 44 in cache entry 3

⇒ Compulsory miss

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
1	0	0x9E768A8A	1
2	2	0x45F36C4D	1
3	1	0xFFB78E14	1
4	3	0xA57B9426	1
5	-	-	0
6	4	0x710F84E3	1
7	-	-	0

3.1- Direct-mapped cache: conflict misses and thrashing

- Conflict miss rate is high
- Very high conflict miss rate = thrashing

Example:

- Load address 44 in cache entry 3
- Load address 77 in cache entry 3, evicts address 44

⇒ Compulsory miss

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
1	0	0x9E768A8A	1
2	2	0x45F36C4D	1
3	2	0x9634E88A	1
4	3	0xA57B9426	1
5	-	-	0
6	4	0x710F84E3	1
7	-	-	0

3.1- Direct-mapped cache: conflict misses and thrashing

- Conflict miss rate is high
- Very high conflict miss rate = thrashing

Example:

- Load address 44 in cache entry 3
- Load address 77 in cache entry 3, evicts address 44
- Load address 44 in cache entry 3, evicts address 76

⇒ Conflict miss

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
1	0	0x9E768A8A	1
2	2	0x45F36C4D	1
3	1	0xFFB78E14	1
4	3	0xA57B9426	1
5	-	-	0
6	4	0x710F84E3	1
7	-	-	0

3.1- Direct-mapped cache: conflict misses and thrashing

- Conflict miss rate is high
- Very high conflict miss rate = thrashing

Example:

- Load address 44 in cache entry 3
- Load address 77 in cache entry 3, evicts address 44
- Load address 44 in cache entry 3, evicts address 76
- Load address 77 in cache entry 3, evicts address 44

⇒ Conflict miss

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
1	0	0x9E768A8A	1
2	2	0x45F36C4D	1
3	2	0x9634E88A	1
4	3	0xA57B9426	1
5	-	-	0
6	4	0x710F84E3	1
7	-	-	0

3.1- Direct-mapped cache: conflict misses and thrashing

- Conflict miss rate is high
- Very high conflict miss rate = thrashing

Example:

- Load address 44 in cache entry 3
- Load address 77 in cache entry 3, evicts address 44
- Load address 44 in cache entry 3, evicts address 76
- Load address 77 in cache entry 3, evicts address 44
- ...

Conflict miss every time from step 3 onward !

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
1	0	0x9E768A8A	1
2	2	0x45F36C4D	1
3	1	0xFFB78E14	1
4	3	0xA57B9426	1
5	-	-	0
6	4	0x710F84E3	1
7	-	-	0

3.2- Fully associative cache

Example:

- Main memory:
 - 256B (addresses: 0 to 255)
- Cache:
 - 32B fully associative
 - 8 cache lines
 - 4B line size
 - 6b tag

	Tag	Data	Valid Bit
0	-	-	0
1	-	-	0
2	-	-	0
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.2- Fully associative cache: loading memory

- CPU request data from main memory with address 155.

$155 \Leftrightarrow 10011011$

- Block size is $4B = 2^2B$, set **2** least significant bits of address to 0 to find block address

$10011011 \& 11111100 = 10011000 \Leftrightarrow 152$

- No need to compute block index, find any empty cache entry
- Use the block address as the tag for the entry

Example:

- Main memory: 256B (addresses: 0 to 255)

address	151	152	153	154	155	156
data	5 B	7 1	0 F	8 4	E 3	9 2

memory block

- Cache: 32B direct-mapped, $8 * 4B$ cache lines

	Tag	Data	Valid Bit
0	152	0x710F84E3	1
1	-	-	0
2	-	-	0
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.2- Fully associative cache: no more conflict misses

- Load address 44 in cache entry 2

⇒ **Compulsory miss**

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	152	0x710F84E3	1
1	44	0xFFB78E14	1
2	-	-	0
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.2- Fully associative cache: no more conflict misses

- Load address 44 in cache entry 2
- Load address 77 in cache entry 3

⇒ Compulsory miss

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	152	0x710F84E3	1
1	44	0xFFB78E14	1
2	76	0x9634E88A	1
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.2- Fully associative cache: no more conflict misses

- Load address 44 in cache entry 2
- Load address 77 in cache entry 3
- Read address 44 in cache entry 2

⇒ Cache hit !

Example:

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	152	0x710F84E3	1
1	44	0xFFB78E14	1
2	76	0x9634E88A	1
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.2- Fully associative cache: no more conflict misses

- Load address 44 in cache entry 2
- Load address 77 in cache entry 3
- Read address 44 in cache entry 2
- Read address 77 in cache entry 3

⇒ Cache hit !

- No more conflict misses, but ...

Example:

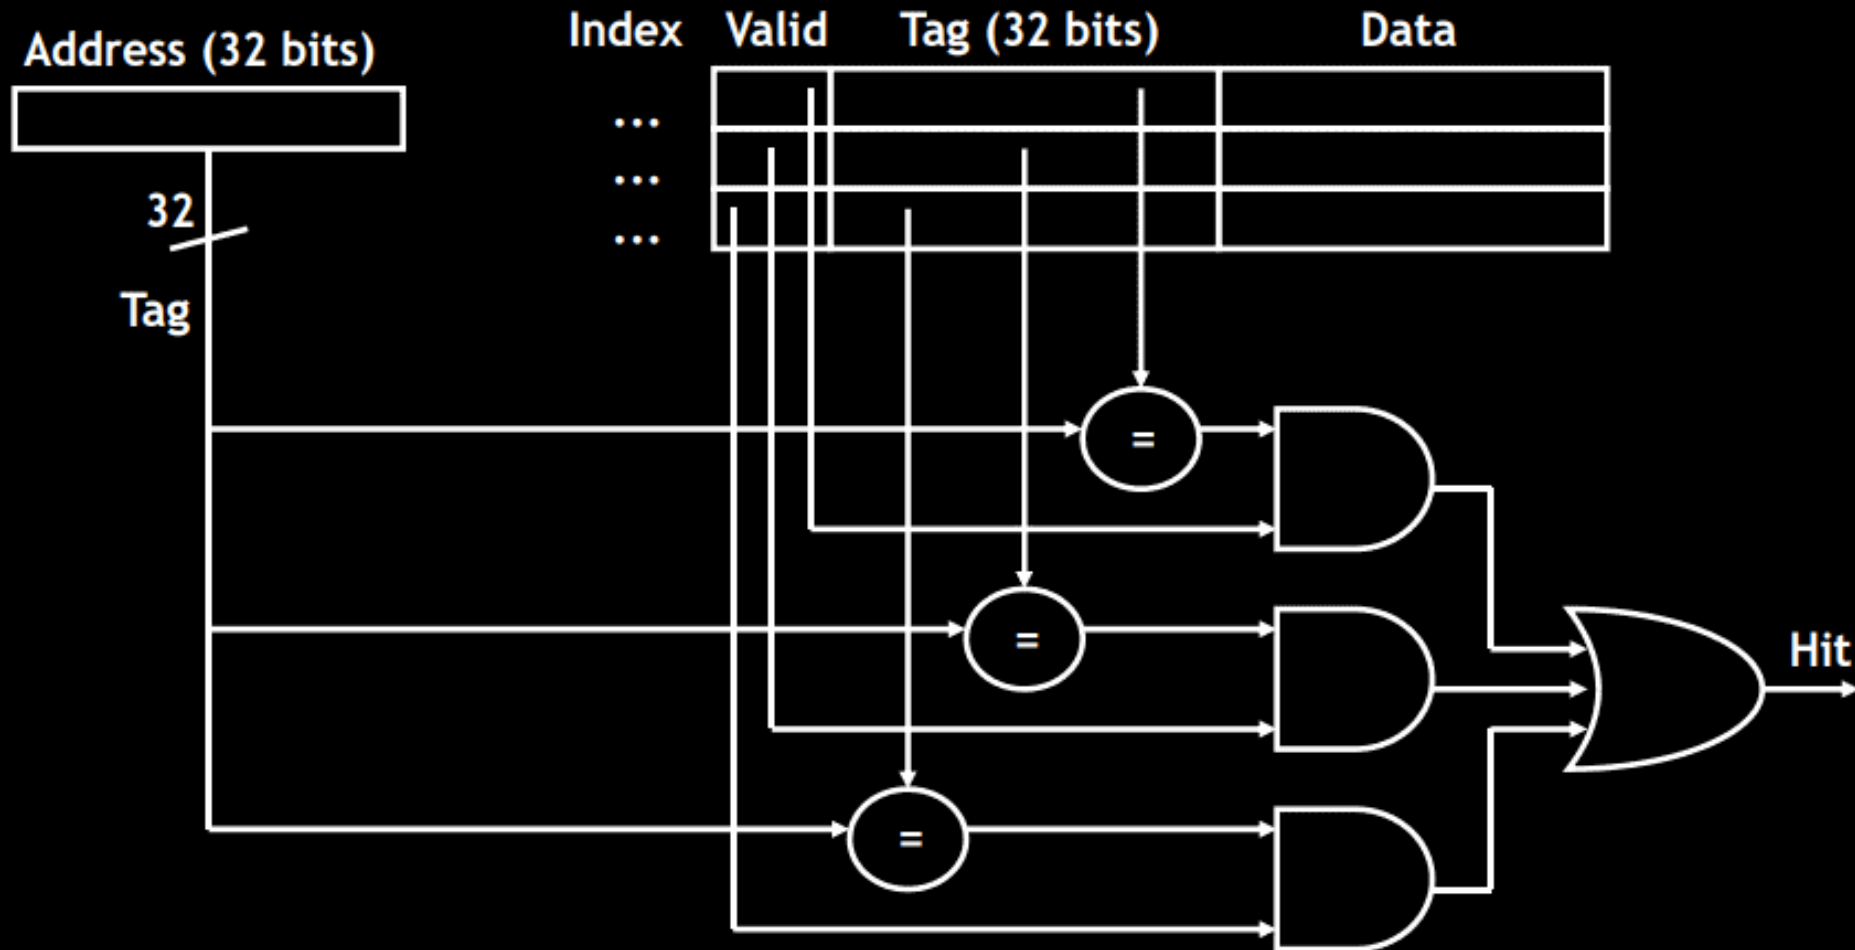
- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	152	0x710F84E3	1
1	44	0xFFB78E14	1
2	76	0x9634E88A	1
3	-	-	0
4	-	-	0
5	-	-	0
6	-	-	0
7	-	-	0

3.2- Fully associative cache: the cost of data request

Example:

- Main memory: 256B (addresses: 0 to 255)
- Cache: 32B direct-mapped, 8 * 4B cache lines



3.2- Set-associative cache

Example:

- Main memory:
 - 256B (addresses: 0 to 255)
- Cache:
 - 32B 2-way associative
 - 4 sets
 - 4B line size
 - 4b tag

	Tag	Data	Valid Bit
0	-	-	0
	-	-	0
1	-	-	0
	-	-	0
2	-	-	0
	-	-	0
3	-	-	0
	-	-	0

3.3- Set-associative cache: loading memory

- CPU request data from main memory with address 155.

155 \Leftrightarrow 10011011

- Block size is 4B = 2^2 B, set 2 least significant bits of address to 0 to find block address

10011011 & 11111100 = 10011000 \Leftrightarrow 152

- Right shift block address by 2 to find block index in main memory (equ. to division by 4)

10011000 \gg 2 = 00100110 \Leftrightarrow 38

Example:

- Main memory: 256B (addresses: 0 to 255)

address	151	152	153	154	155	156
data	5 B	7 1	0 F	8 4	E 3	9 2

memory block

- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
	-	-	0
1	-	-	0
	-	-	0
2	-	-	0
	-	-	0
3	-	-	0
	-	-	0

3.3- Set-associative cache: loading memory

- There are $4 = 2^2$ cache sets, keep the **2** least significant bits of that index to find the set index (equ. to mod 4)

$$00100110 \& 00000011 = 00000010 \Leftrightarrow 2$$

- Find an empty cache entry in the set
- Use the **remaining** bits as the tag.

$$(00100110 \& 11111100) \gg 2 = 00001001 \Leftrightarrow 9$$

Example:

- Main memory: 256B (addresses: 0 to 255)

address	151	152	153	154	155	156
data	5 B	7 1	0 F	8 4	E 3	9 2

memory block

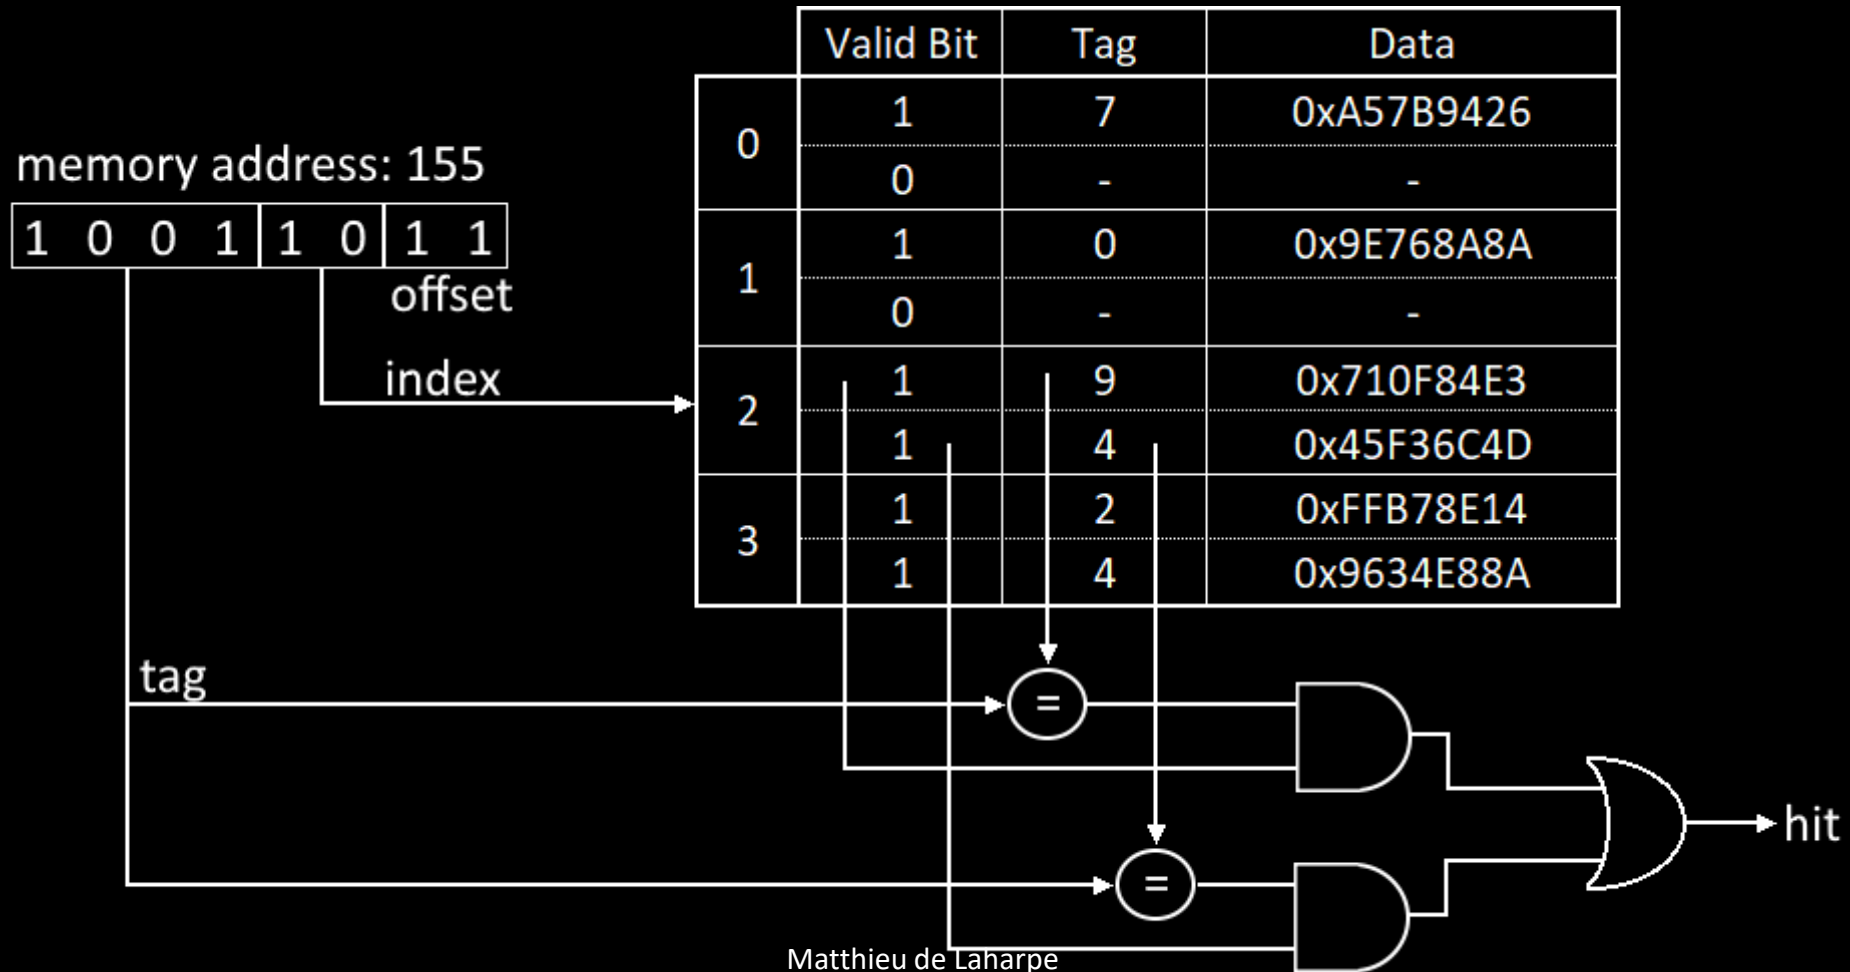
- Cache: 32B direct-mapped, 8 * 4B cache lines

	Tag	Data	Valid Bit
0	-	-	0
	-	-	0
1	-	-	0
	-	-	0
2	9	0x710F84E3	1
	-	-	0
3	-	-	0
	-	-	0

3.3- Set-associative cache: balanced cost for data requests

Example:

- Main memory: 256B (addresses: 0 to 255)
- Cache: 32B direct-mapped, 8 * 4B cache lines



4- Cache policies


- CPU cache is a critical component regarding performances

L1 cache ref	0,5 ns	
L2 cache ref	7 ns	14 * L1 cache
Main memory ref	100 ns	14 * L2 cache, 200 * L1 cache

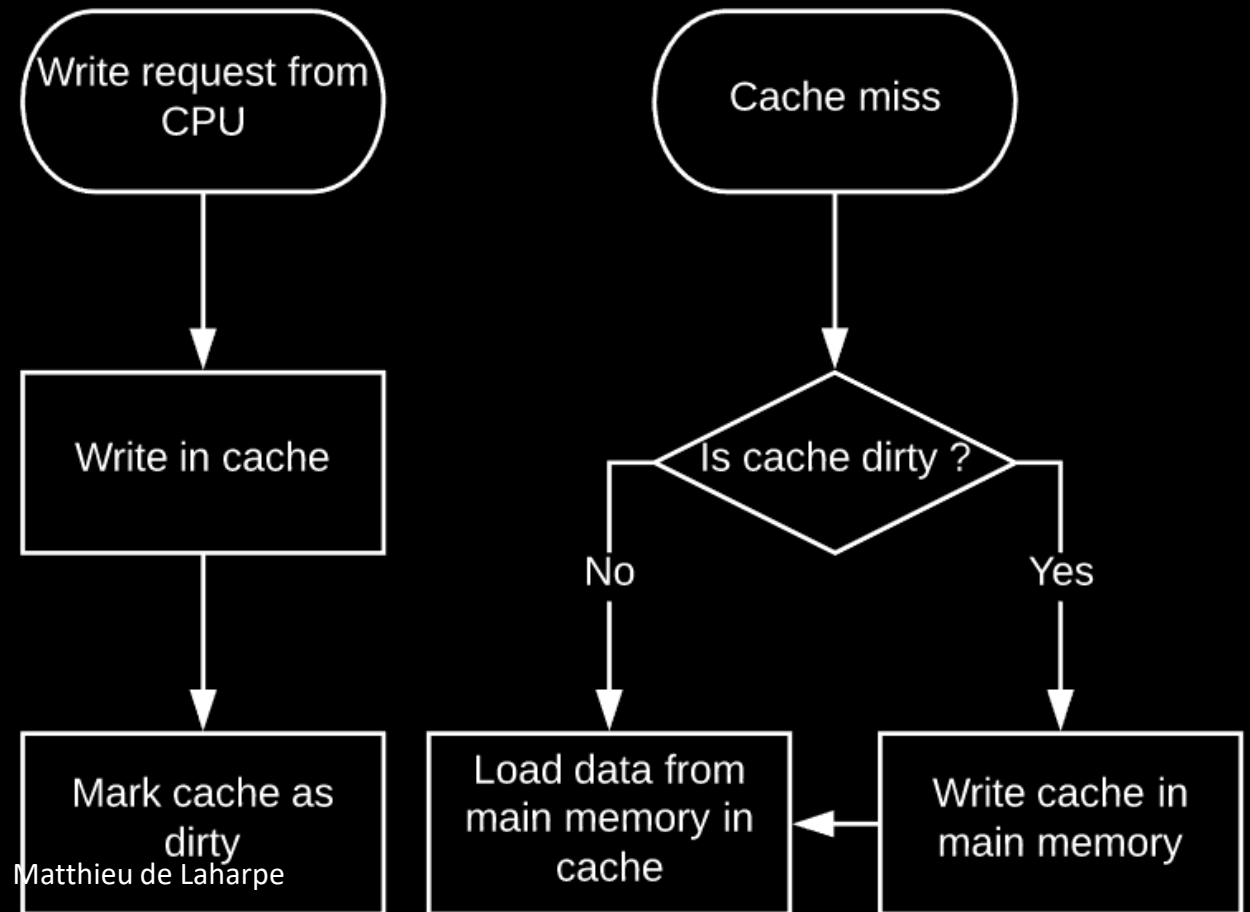
- CPU cache behaves differently depending on its characteristics, environment and on the applications that use it
- A lot of different policies have been defined to dictate cache behavior and adapt to different situations

4.1- Write policies:

Write through



Write back (requires a dirty flag in the cache entry)



4.1- Write policies

When a write triggers a cache miss:

- Write allocate: load data at the missed location to the cache
- No-write allocate: write directly to the main memory/next level cache

Policies are usually paired in this way:

- Write back + write allocate -> hoping for subsequent w/r to the same location
- Write through + no write allocate

Intermediate write through policy:

- Write operations are buffered together to save time, power and bus utilization
- AMD Bulldozer architecture uses a specialised Write Coalescing Cache between L1d and L2


4.2- Replacement policies

Which cache entry should be evicted to make room for a new cache line ?

- Wide variety of policies
 - Least Recently Used (LRU) -> temporal locality, very easy on 2-way associative cache (1bit)
 - Most Recently Used (MRU) -> useful when sequentially looping over a data set
 - First In first Out (FIFO) -> simplest LRU
 - Least Frequently Used (LFU) -> can lead to cache pollution
 - Random
 - ...
- A lot of implementations are variations, improvements or hybrids of LRU
 - LRFU, EELRU, LRU-K, LRU-2 ...
 - Tree Pseudo LRU, bit Pseudo LRU ...

4.2- Replacement policies: Pseudo LRU tree

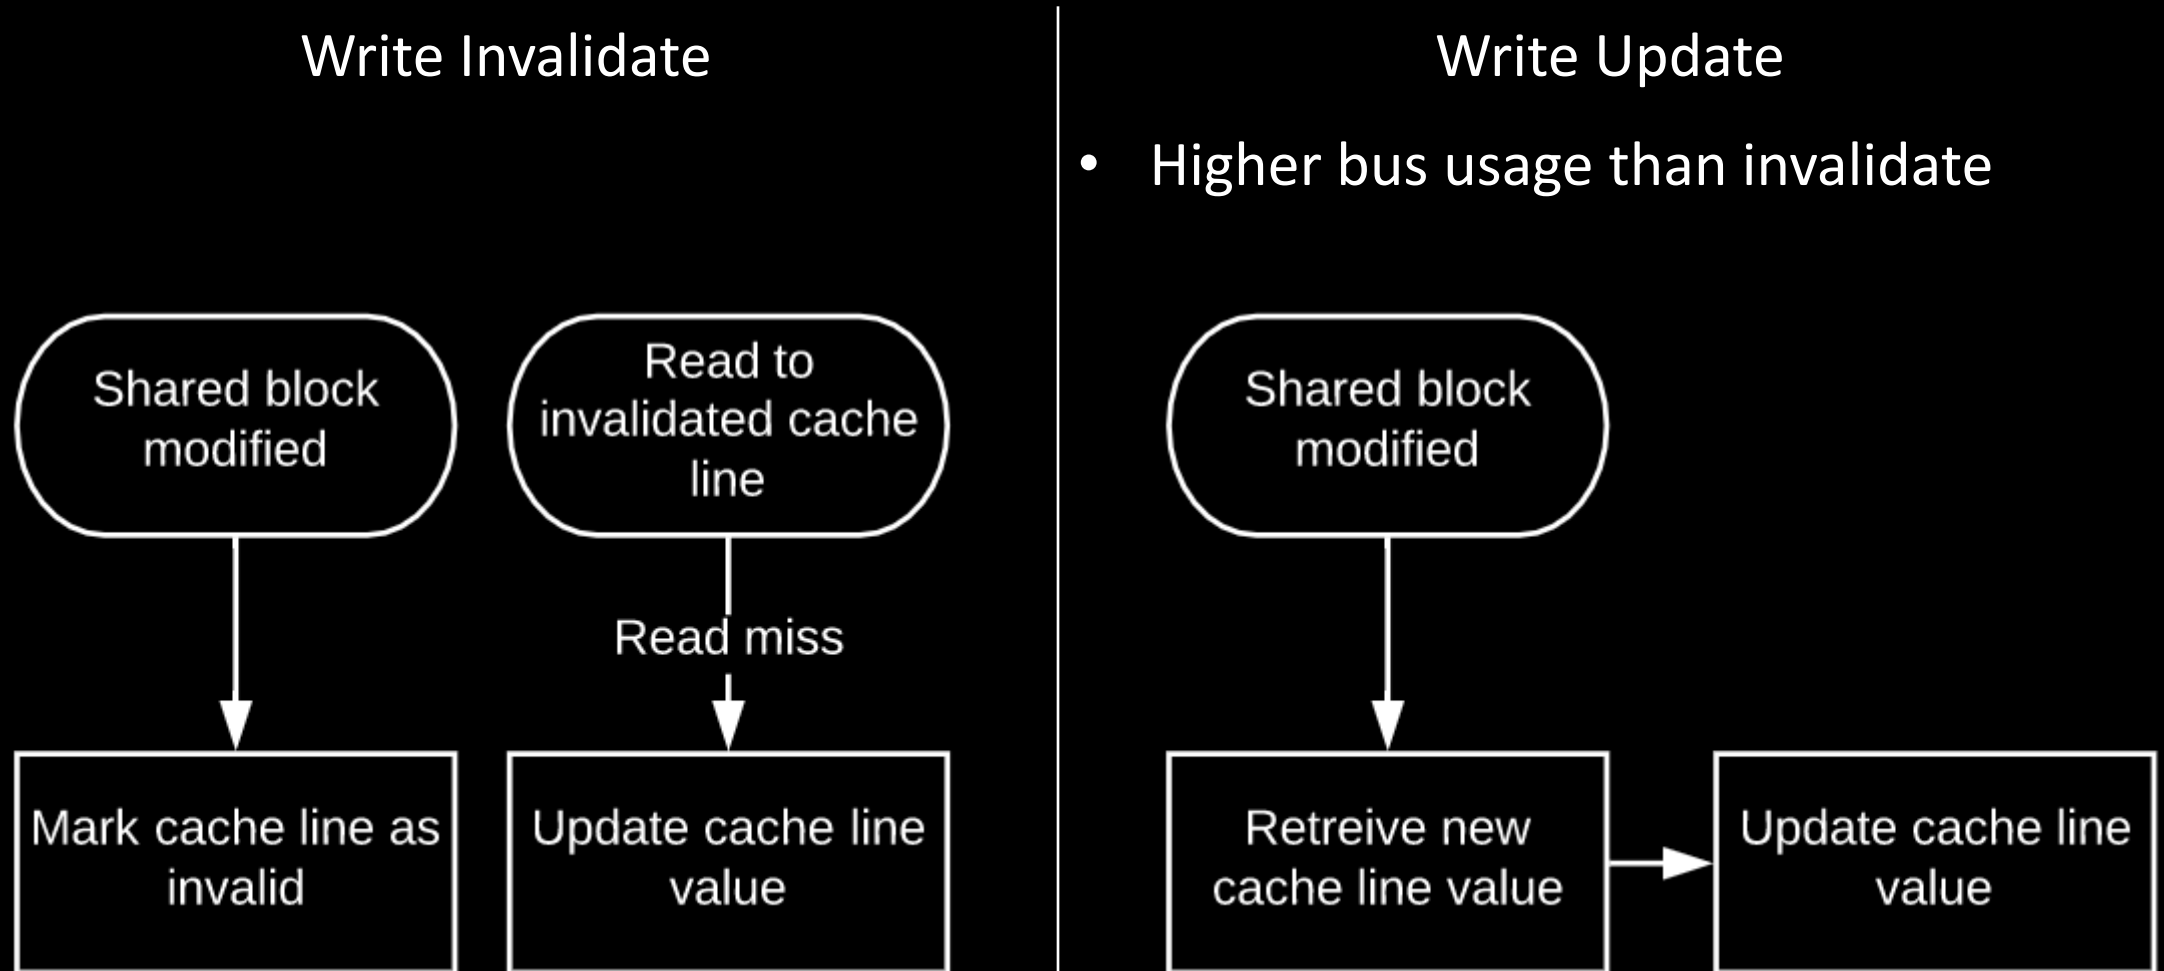
- Requires N-1 bits per set for a N-way set-associative cache
 - On a 256KiB 4-way associative cache with a 64B line size: (L2 cache in Intel Kaby Lake μ arch)
 $256\text{KiB} / (4 * 64\text{B}) = 1024$ sets
 $3\text{b} * 1024 = 3\text{Kib} = 384\text{B} = 0,375\text{KiB} (\approx 0,15\% \text{ of cache size})$



5- Cache coherence

- Maintain uniformity of shared data resources across multiple local caches in multi CPU system
- When a CPU modifies the value of a shared data, this modification must be propagated to all the other caches
- Two most common mechanisms:
 - Bus snooping: cache have a coherency controller that monitors the bus for changes in shared memory blocks
 - Directory-based

5- Cache coherence: bus snooping



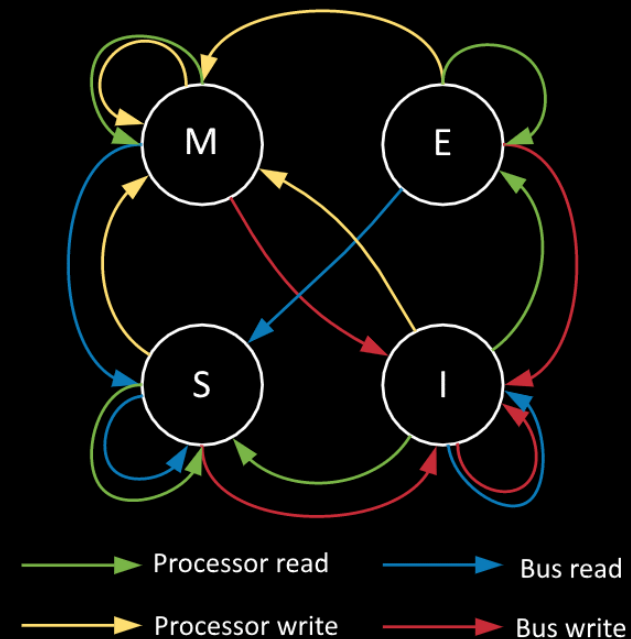
5- Cache coherence: MESI protocol

- Cache line can be in 4 different states:
 - Modified: cache line is present in 1 cache, doesn't match the backing store
 - Exclusive: cache line is present in 1 cache, matches the backing store
 - Shared: cache line is present in multiple caches, matches the backing store
 - Invalid: cache line is invalid

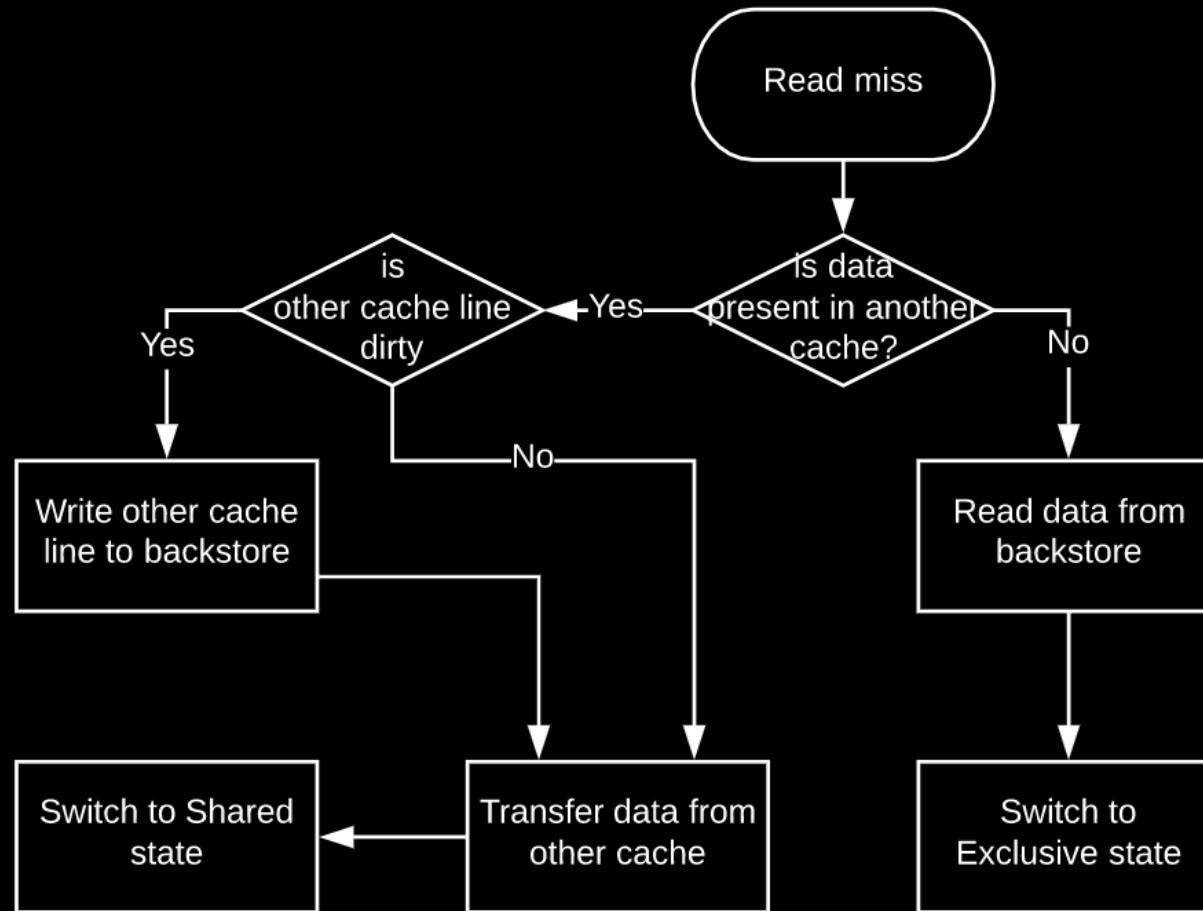
Relative state of cache lines
with the MESI protocol

	M	E	S	I
M	✗	✗	✗	✓
E	✗	✗	✗	✓
S	✗	✗	✓	✓
I	✓	✓	✓	✓

MESI state
transitions



5- Cache coherence: MESI protocol



6- Going further

- Address translation
 - Modern systems use virtual memory
 - Need translation between virtual and physical addresses
 - Specialized caches handle address translations: Translation Lookaside Buffer
- Other specialized caches
 - Victim cache: holds evicted cache lines to reduce conflict misses
 - Trace cache, micro-operation cache ...
- Cache pollution
- Cache control instructions
- Branch prediction and cache prefetching

References

Wikipedia

- [Cache \(computing\)](#)
- [CPU cache](#)
- [Cache coherence](#)
- [Cache replacement policies](#)
- [Pseudo LRU](#)
- [Cache pollution](#)
- [Cache performance](#)
- [Thrashing \(computer science\)](#)
- [Cache prefetching](#)
- [Bus snooping](#)
- [MESI protocol](#)

University of Washington, Computer Science & Engineering [lectures from winter 2011](#):

- Lecture 16: [Introduction to caches](#)
- Lecture 17: [Locality & Cache organization](#)
- Lecture 18: [More on cache organization](#)
- Lecture 19-20: [Cache writing](#)

[wikichip.org: i7-7700HQ](http://wikichip.org:i7-7700HQ)

[wikichip.org: Kaby Lake micro arch Memory Hierarchy](http://wikichip.org:Kaby Lake micro arch Memory Hierarchy)

[Northeastern University: cache basics](#)

[Carnegie Mellon University: cache associativity](#)

[University of Wisconsin-Madison: Replacement policies](#)

[lwn.net: CPU Caches](#)

MIT Computer Science and AI Lab:

- [Adaptive Insertion Policies for High Performance Caching](#)
- [Modeling Cache Performance Beyond LRU](#)

[General Adaptive Replacement Policies](#)

[Latency numbers every programmer should know](#)

[The memory wall fallacy](#)